

CS6611- CREATIVE AND INNOVATIVE PROJECT

B.E CSE VI Q - BATCH

TEAM MEMBERS:

TEAM NO: 12

S.NO.	REG. NO.	NAME
1.	2018103569	Niranjana K
2.	2018103063	Shankar Kumar S
3.	2018103620	Vedanth S

Project Title: Identifying Sound Labels from Spectrograms using Deep Neural Networks

WEEK 11 - EXECUTION (90% IMPLEMENTATION)

Observation document (5)	
On the Spot exercise (5)	
Laboratory exercises identified (15)	
Total (25)	

TRAINING A MODEL ON MFCC-SPECTROGRAM

```
In [5]: x_train_mfcc = x_train_mfcc[1:,:,:]
```

```
In [6]: x_train_mfcc = x_train_mfcc[:, :, :, np.newaxis]
```

```
In [7]: y_train_mfcc = []
        for i in range(3799):
            y_train_mfcc.append(train_mfcc_spec_normalised[i][1])
```

```
In [8]: y_train_mfcc = np.array(y_train_mfcc)
```

ARCHITECTURE

```
In [9]: def create_keras_model(model_input, num_classes):
        x = Conv2D(64, kernel_size=3, activation="relu",padding='same')(model_input)
        x = MaxPooling2D(pool_size=(2, 2))(x)
        x = Conv2D(128, kernel_size=3,padding='same', activation="relu")(x)
        x = MaxPooling2D(pool_size=(2, 2))(x)
        x = Conv2D(256, kernel_size=3,padding='same', activation="relu")(x)
        x = MaxPooling2D(pool_size=(2, 2))(x)
        x = Conv2D(256, kernel_size=3,padding='same', activation="relu")(x)
        x = MaxPooling2D(pool_size=(2, 2))(x)
        x = Flatten()(x)
        x = Dense(256, activation = 'relu')(x)
        x = Dropout(0.5)(x)
        x = Dense(num_classes, activation="softmax")(x)
        model = Model(model_input, x, name='model_1')
        return model
```

```
In [13]:
callbacks = [
    tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10, verbose=1,
    mode="min"),
    tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss',
                                         mode='min',
                                         verbose=1,
                                         patience=5,
                                         min_delta=0.0001,
                                         factor=0.2
                                         ),
]
```

MODEL SUMMARY:

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 20, 216, 1)]	0

conv2d (Conv2D)	(None, 20, 216, 64)	640

max_pooling2d (MaxPooling2D)	(None, 10, 108, 64)	0

conv2d_1 (Conv2D)	(None, 10, 108, 128)	73856

max_pooling2d_1 (MaxPooling2D)	(None, 5, 54, 128)	0

conv2d_2 (Conv2D)	(None, 5, 54, 256)	295168

max_pooling2d_2 (MaxPooling2D)	(None, 2, 27, 256)	0

conv2d_3 (Conv2D)	(None, 2, 27, 256)	590080

```

-----
max_pooling2d_3 (MaxPooling2 (None, 1, 13, 256)      0
-----
flatten (Flatten)          (None, 3328)              0
-----
dense (Dense)              (None, 256)              852224
-----
dropout (Dropout)          (None, 256)              0
-----
dense_1 (Dense)            (None, 50)              12850
=====
Total params: 1,824,818
Trainable params: 1,824,818
Non-trainable params: 0
-----

```

COMPILING THE MODEL:

```

In [15]: model_mfcc_1.compile(optimizer='adam', loss='sparse_categorical_crossentropy', met
         rics=['accuracy'])

```

```

In [16]: # Training and Evaluation of the model
         h1 = model_mfcc_1.fit(x_train_mfcc, y_train_mfcc, verbose=1, batch_size = 64 , epo
         chs=200, validation_split=0.1, callbacks = callbacks)

```

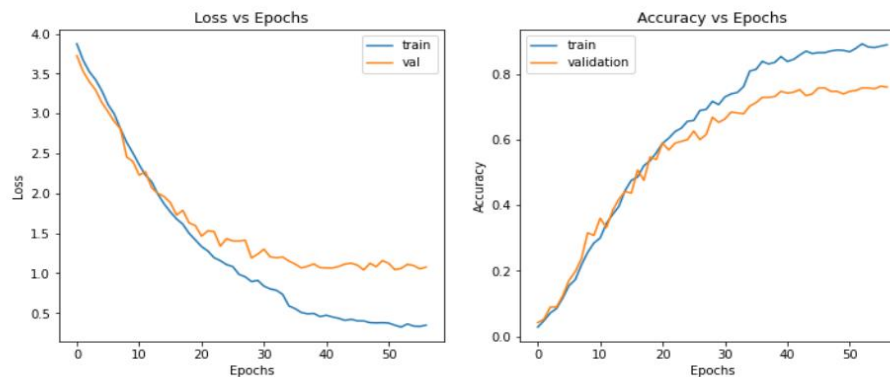
TEST RESULTS:

```
Epoch 53/200
54/54 [=====] - 1s 16ms/step - loss: 0.3368 - accuracy: 0.8897 - val_loss: 1.0606 - val_accuracy: 0.7579
Epoch 54/200
54/54 [=====] - 1s 16ms/step - loss: 0.3579 - accuracy: 0.8916 - val_loss: 1.1111 - val_accuracy: 0.7579
Epoch 55/200
54/54 [=====] - 1s 16ms/step - loss: 0.3314 - accuracy: 0.8799 - val_loss: 1.0971 - val_accuracy: 0.7553
Epoch 56/200
54/54 [=====] - 1s 16ms/step - loss: 0.3328 - accuracy: 0.8906 - val_loss: 1.0553 - val_accuracy: 0.7632
Epoch 57/200
54/54 [=====] - 1s 16ms/step - loss: 0.3530 - accuracy: 0.8843 - val_loss: 1.0778 - val_accuracy: 0.7605

Epoch 00057: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.
Epoch 00057: early stopping
```

LOSS AND ACCURACY GRAPHS

```
In [18]: plot_metrics(h1)
```



INPUT: 3800X20X216 MFCC spectrogram images

OUTPUT: A model with learned weights

ARCHITECTURE:

```
x = Conv2D(64, kernel_size=(3, 3), activation='relu',padding='same')(model_input)
x = BatchNormalization()(x)

x = Conv2D(128, kernel_size=(3, 3), activation='relu',padding='same')(x)
x = BatchNormalization()(x)

x = Conv2D(128, kernel_size=(3, 3), activation='relu',padding='same')(x)
x = BatchNormalization()(x)

x = MaxPooling2D(pool_size=(4, 4))(x)
x = Dropout(0.25)(x)

x = Flatten()(x)

x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.25)(x)

x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.25)(x)

x = Dense(64, activation='relu')(x)

x = BatchNormalization()(x)
x = Dropout(0.25)(x)

x = Dense(num_classes, activation='softmax')(x)
model = Model(model_input,x,name = 'model_2')
return model
```

COMPILING THE MODEL

```
In [21]: model_mfcc_2.compile(loss='sparse_categorical_crossentropy', optimizer=tf.keras.  
optimizers.Adam(), metrics=['accuracy'])
```

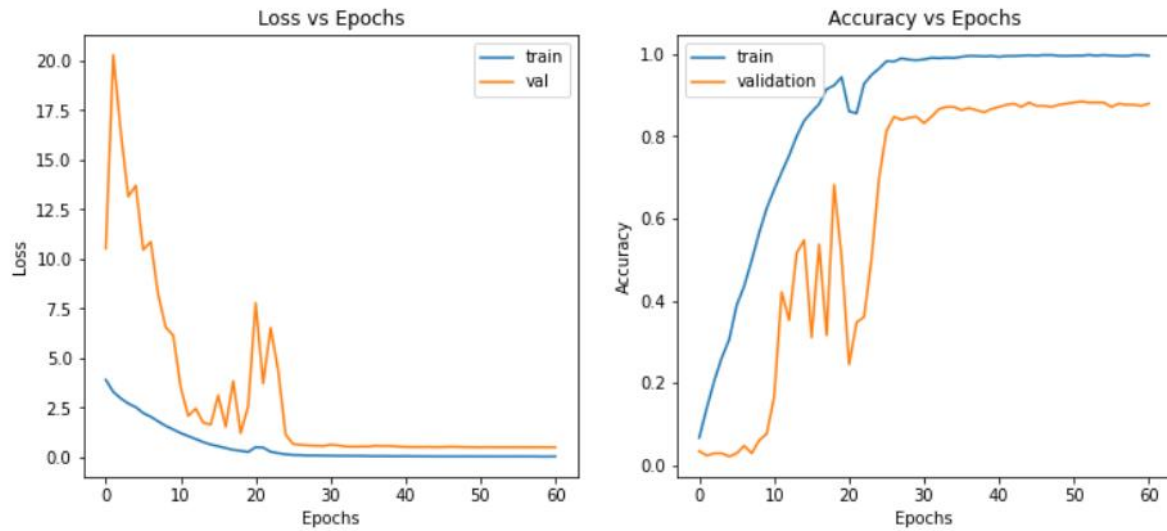
```
In [22]: h2 = model_mfcc_2.fit(x_train_mfcc, y_train_mfcc, batch_size=64, epochs=200, ver  
bose=1, validation_split=0.1, callbacks = callbacks)
```

RESULTS

```
Epoch 60/200  
54/54 [=====] - 3s 48ms/step - loss: 0.0371 - accuracy: 0.9972 - val_loss: 0.5062 - val_accuracy: 0.8737  
Epoch 61/200  
54/54 [=====] - 3s 47ms/step - loss: 0.0411 - accuracy: 0.9952 - val_loss: 0.5065 - val_accuracy: 0.8789  
  
Epoch 00061: ReduceLRonPlateau reducing learning rate to 1.6000001778593287e-06.  
Epoch 00061: early stopping
```

LOSS & ACCURACY GRAPHS

```
]:\nplot_metrics(h2)
```



INPUT: 3800X20X216 MFCC spectrogram images

OUTPUT: A model with learned weights